

FINDING

Maldros ran a full analytics investigation — from a raw business question to an executive briefing — autonomously in 288.9 seconds. Governance was never traded for that speed: the L1 gate blocked that run's own first draft, blocked 14 of 18 Storyteller runs across Phase 1 development, and released a report only once every claim traced to a source — with a human sign-off gate that never auto-approved.

IMPACT

The speed is the least interesting part. The differentiator is that it cost nothing in accountability — governance enforced mechanically, as hard rules the language model cannot talk its way past, is the thread this case study returns to at every layer. The real finding was one aggregate monitoring structurally misses: the Q1 population abuse rate was not elevated — it sat at 0.946× the non-Q1 average — while a 41-account cluster ran at 16.54× baseline, inside a quarter whose company-wide API-abuse impact was \$5.89M (synthetic domain) — a concentration no dashboard watching the aggregate would surface.

Everything here is validated via simulation against a synthetic dataset; nothing is production-deployed.

EVIDENCE · THE VALIDATED RUN

Autonomous end-to-end run · run 6cf9b15c 288.9s Wall-clock for the full Orchestrator → Analyst → Statistician → Storyteller pipeline, against an estimated 2–5 analyst-days by hand. Runtime only — it carries no governance claim of its own. Validated via simulation.	Governance held · the approved report 97.9% 47 of 48 claims cited on the report approved at the Confirmation Gate (Discovery Report 1b4f50b2, run ff0b6f70, AIMS Mode B f3d5a232), the run that cleared all five deterministic L1 checks. 14 of 18 Storyteller runs were blocked across Phase 1 development. No auto-approval at the Confirmation Gate.	The signal every aggregate missed 16.54× A 41-account cluster, while the Q1 population rate sat at 0.946×; company-wide Q1 API-abuse impact: \$5.89M (synthetic domain).	Capability phases validated 6 Investigation, self-healing, experiment analysis, self-improvement, invention, and full autonomous audit — each with a named artifact trail.
--	---	--	--

EVIDENCE · THE ECONOMICS AT CONCURRENT FRONTIER SCALE

Incremental cash cost · per year \$210K–\$800K Across 15–60 analyst seats (~750–15,000 investigations/yr). Dominated by platform-engineering headcount; tokens and warehouse under 10%. [Projected]	Return on the general composite ~45:1–50:1 Scaled Evidence-8 value over Layer-A cash cost; ~84:1–98:1 including the fraud-domain component. [Projected]	Cost per investigation \$49–\$300 Falls as concurrent volume rises, against a \$26,250 traditional human investigation unit. [Projected]	Payback on one-time integration ~4–5 days One-time production-integration cost against a far larger monthly net value. Arithmetically correct, not a claim of instant value — plan to the milestone sequence. [Projected]
---	---	--	---

The signal was concentration, not volume.

A 41-account cluster ran at 16.54× the baseline abuse rate while the Q1 population rate sat at 0.946× — slightly below normal. No monitor watching the aggregate would ever surface it; the signal lives entirely in the concentration.




Figure. Phase-1 generated finding, validated via simulation — the real output of the run, not a screenshot of a live production system. Rates are expressed relative to the baseline abuse rate (1.0x).

RECOMMENDATION

What a reader should do with this.

For a technical evaluator: load Discovery Report 1b4f50b2 and AIMS Mode B briefing f3d5a232 and verify the citation trail directly — the 97.9% figure is not a self-reported claim but the output of a deterministic citation-coverage check that blocked 14 of 18 Storyteller runs over the course of Phase 1 development. Confirm from the discovery path alone that the 16.54× cluster and the 0.946× population rate reproduce — the system holds itself to that standard, so hold it there too.

For a decision-maker: treat the 288.9-second runtime as the least important fact here. The governance-intact result — 97.9% citation coverage, 47 of 48 claims, on the report approved at the Confirmation Gate under five deterministic L1 checks with no auto-approval — is what separates this build from a faster, less careful alternative. At concurrent frontier scale the economics follow: a ~45:1 general-composite ROI on a ~\$210K–\$800K incremental cash cost, best budgeted as a platform-engineering line, not a compute line — the binding cost is the people and one-time integration to operate the system, not marginal compute.

Maturity states, stated plainly. Maldros is a single-operator simulation built and run via scripted agent-session execution against a synthetic dataset. Capabilities are labeled throughout as design-complete, validated via simulation, or production-deployed — and nothing here qualifies as the last. The engineering rigor, not any claimed production traction, is the differentiator.